T18 # Models & Thinking Modes

Technical Beginner

Choosing the right model and the right level of reasoning

Current models

Standard API rates verified October 27, 2026. USD per million input/output tokens, excluding caching.

Model	Input / output	Context	Default effort
Haiku 5.5	$0.10 / $0.50 (over 100K: $0.50 / $2.50)	1M	medium
Sonnet 5.5	$2 / $10	1M	medium
Opus 5.5	$4 / $20	1M	medium
Fable 5.1	$10 / $50	1M	high

Sonnet's published rate remains $2/$10 after the launch promotion. Opus 5.5 is the direct-service default; account policy can override it. Model choice does not establish a typical task cost. Effort shown is the Claude Code default; the API default for Sonnet 5.5 is `high`.

Select and save

`/model` opens the picker. **Enter** saves a default; **s** applies it to this session only. `claude --model sonnet` applies to the launched session. Project and managed settings can reapply at the next launch.

Aliases depend on the provider. `opus` maps to 5.5 on most providers, but 4.6 on Foundry. `sonnet` maps to 5.5 on the direct API, 4.6 on Claude Platform on AWS, and 4.5 on Bedrock, Google, and Foundry.

`haiku` selects 5.5 on the Anthropic API (v2.1.293+), 4.5 on AWS, Bedrock, Google, and Foundry. `fable` selects 5.1 where available, or 5 through the Claude apps gateway. `best` selects Fable when available, otherwise Opus. Pin a supported full ID when version identity matters.

Thinking and effort

Opus 5.5, Sonnet 5.5, Haiku 5.5, and Fable 5.1 use adaptive thinking. Legacy Haiku 4.5 has **no effort parameter**.

Current adaptive models support `low`, `medium`, `high`, `xhigh`, and `max`. Opus 4.6 and Sonnet 4.6 support `max` but not `xhigh`.

`/effort` saves levels through `xhigh` per model. Use **s** for a session-only choice. `max` is session-only; `/effort auto` clears the active model's saved choice.

`Option+T` / `Alt+T` toggles thinking where supported. It has no effect on Opus 5.5, Sonnet 5.5, Haiku 5.5, or Fable, whose thinking cannot be turned off. Lower effort for bounded work instead.

`ultrathink` adds a reasoning instruction without changing API effort. `think hard` is ordinary prompt text.

Choose by measured results

Evaluate Haiku on bounded tasks, Sonnet on routine work, and Opus or Fable on harder tasks. Compare accepted outcomes, retries, latency, and total cost on the same task set. Teammates can use different models.

Sources: Model configuration, API pricing.