T19 Context Window: 200K vs 1M

Technical Intermediate

When to switch to the extended context window and at what cost

Context is model-specific

Direct API model	Context	Maximum output
Haiku 5.5	1M native	128K
Haiku 4.5 (legacy)	200K	64K
Sonnet 5.5	1M native	128K
Opus 5.5	1M native	128K
Fable 5.1	1M native	128K

Guide model snapshot, October 27, 2026. Provider settings can expose different limits. Inspect `/context` for the active session.

No native-1M premium

Opus, Sonnet, and Fable 5.x use standard rates beyond 200K input tokens (Haiku 5.5 is priced higher above 100K). They do not need the old long-context beta header. Sonnet 5.5 has no 200K variant on the direct API.

Older Opus 4.6 and Sonnet 4.6 have `[1m]` variants with plan/provider restrictions. Do not transfer their access rules to current models.

Price a complete request

Illustrative USD totals without caching or retries:

Input / output	Sonnet 5.5	Opus 5.5
50K / 5K	$0.15	$0.30
150K / 20K	$0.50	$1.00
500K / 50K	$1.50	$3.00

Use separate input, output, cache-read, and cache-write counts. A context-window size is capacity, not a bill or an accuracy score.

Control accumulation

`/context` shows the token breakdown. `/compact` summarizes the conversation; `/clear` starts fresh. Save decisions and verification commands before resetting context.

`/autocompact 500k` saves a smaller compaction window. `/autocompact auto` restores the model default. Native 1M sessions normally compact around 967K.

Load relevant files rather than filling the available window. Split independent tasks when separate contexts help. No fixed percentage guarantees reliable recall. Idle compaction (v2.1.299) compacts a 200K+ conversation idle about 54 minutes before its one-hour cache expires; `"idleCompaction": false` disables it.

Evaluate retrieval

Older MRCR results for Opus 4.6 or Sonnet 4.5 remain results for those models and benchmark conditions. They do not establish retrieval quality for Sonnet 5.5 or Opus 5.5.

Test on your own document set with answerable questions, expected citations, and distractors. Measure missed evidence as well as correct answers.

Sources: Model configuration, model specifications, pricing.