T21 Fast Mode & API Breaking Changes

Technical Intermediate

Fast mode and the major API changes you need to know

Fast mode on Opus 5.5

Fast mode uses faster serving of the same model. Anthropic advertises **up to 2.5 times faster output**, not a guarantee for every request. Claude Code selects Opus 5.5 for fast mode from v2.1.280.

Rate per million tokens	Standard	Fast
Input	$4	$8
Output	$20	$40

Opus 5 and Opus 4.8 also support fast mode at $10/$50. Sonnet and Haiku do not. Opus 4.7 is no longer supported in fast mode.

Activation and billing

Use `/fast` or `Option+O` / `Alt+O`. The choice normally persists between sessions; organization settings and availability still apply. Subscription fast mode uses usage credits, not the included subscription allowance.

Compare observed latency and cost on the same task. Faster output does not remove tool waits, retries, or review work. Sonnet 5.5 standard rates are $2/$10, four times lower than Opus 5.5 fast rates for the same token counts.

API integration

Use the current API fast-mode instructions for supported model IDs, headers, request fields, and fallback behavior. Do not copy an older beta example without checking SDK compatibility.

For adaptive reasoning on current models:

```
response = client.messages.create(
    model="claude-opus-5-5",
    max_tokens=16000,
    thinking={"type": "adaptive"},
    output_config={"effort": "high"},
    messages=[{"role": "user", "content": "Review this design."}],
)
```

This is a standard-speed example. Opus 5.5 defaults to medium effort in Claude Code; the example explicitly asks for high.

Migrating older examples

Opus 4.6 removed assistant prefill. For schema-constrained JSON, use `output_config.format`; for strict tool inputs, put `strict: true` on the tool definition. A system instruction alone does not guarantee valid JSON.

Current Opus, Sonnet 5.5, Haiku 5.5, and Fable use adaptive thinking. Fixed `budget_tokens` examples belong to older models.

Legacy Haiku 4.5 has no effort parameter.

Verified September 24, 2026. Sources: Claude Code fast mode, structured outputs.